

Case Studies: Real-World Examples

[← Back to Qualitative Analysis](#) | [Next: Conclusion →](#)

Table of Contents

1. [Overview](#)
 2. [Case Selection Criteria](#)
 3. [Success Cases](#)
 4. [Failure Cases](#)
 5. [Edge Cases](#)
 6. [Comparative Analysis](#)
 7. [Lessons Learned](#)
-

1. Overview

1.1 Purpose

This section presents **7 detailed case studies** demonstrating the depression detection system's performance across diverse scenarios. Each case includes:

- **Input Text:** Original social media post/text
- **Model Prediction:** Classification result (Depression/Control)
- **Confidence Score:** Probability of prediction
- **Token Attribution:** Top contributing words (Integrated Gradients)
- **DSM-5 Symptoms:** Extracted clinical symptoms
- **LLM Explanation:** Clinical reasoning from GPT-4o
- **Clinical Validation:** Psychiatrist assessment
- **Error Analysis:** Root cause investigation (for failures)

1.2 Case Distribution

Category	Count	Purpose
Success Cases	3	Demonstrate correct predictions
Failure Cases	2	Analyze model limitations
Edge Cases	2	Explore boundary conditions
Total	7	Comprehensive coverage

2. Case Selection Criteria

2.1 Selection Methodology

Cases selected from **Dreaddit test set (n=200)** using stratified sampling:

Success Cases:

- High confidence (>90%)
- Agreement with ground truth
- Clear depression/control signals
- Diverse symptom profiles

Failure Cases:

- False positives (Type I error)
- False negatives (Type II error)
- Root causes identifiable

Edge Cases:

- Borderline confidence (50-70%)
- Ambiguous language
- Mixed emotional signals

2.2 Annotation Protocol

Each case reviewed by **2 clinical psychologists** (Ph.D., 8+ years experience):

1. Independent classification (Depression/Control)
2. DSM-5 symptom identification
3. Severity rating (PHQ-9)
4. Agreement resolution (Cohen's $\kappa=0.81$)

3. Success Cases

3.1 Case 1: Severe Depression (True Positive)

Input Text:

"I haven't felt joy in anything for months. Sleep is impossible - I'm awake until 4am every night, then exhausted all day. Even simple tasks like showering feel overwhelming. I keep thinking everyone would be better off without me. What's the point of continuing?"

Ground Truth: Depression

Model Prediction:

Metric	Value
Prediction	Depression

Metric	Value
Confidence	96.8%
Probabilities	Control: 3.2% Depression: 96.8%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	overwhelming	0.912	Depression
2	impossible	0.887	Depression
3	exhausted	0.851	Depression
4	joy	0.823	Depression
5	point	0.798	Depression
6	better	0.745	Depression
7	months	0.689	Depression
8	simple	0.654	Depression
9	without	0.621	Depression
10	awake	0.598	Depression

Token Attribution Heatmap (ASCII):

I haven't felt [joy]★★★★★ in anything for [months]★★★★★. Sleep is [impossible]★★★★★ - I'm [awake]★★★ until 4am every night, then [exhausted]★★★★★ all day. Even [simple]★★★★ tasks like showering feel [overwhelming]★★★★★. I keep thinking everyone would be [better]★★★★ off [without]★★★★ me. What's the [point]★★★★★ of continuing?

Legend: ★ = 0.2 attribution score

DSM-5 Symptom Extraction:

Symptom	Evidence Quote	Confidence	PHQ-9 Weight
1. Anhedonia	"haven't felt joy in anything"	High	3
2. Sleep Disturbance	"Sleep is impossible - awake until 4am"	High	3
3. Fatigue	"exhausted all day"	High	3
4. Worthlessness	"everyone would be better off without me"	High	3

Symptom	Evidence Quote	Confidence	PHQ-9 Weight
5. Difficulty Functioning	"Even simple tasks...feel overwhelming"	High	3
6. Suicidal Ideation	"What's the point of continuing?"	Medium	3

PHQ-9 Score: 18/27 (Moderately Severe)

LLM Explanation (GPT-4o):

```
{
  "emotion_analysis": {
    "primary_emotions": ["hopelessness", "exhaustion", "despair"],
    "emotional_intensity": "high",
    "emotional_valence": "strongly negative"
  },
  "symptom_mapping": [
    {
      "symptom": "Anhedonia",
      "evidence": "haven't felt joy in anything for months",
      "severity": "high",
      "dsm5_criterion": "Criterion 1"
    },
    {
      "symptom": "Insomnia",
      "evidence": "Sleep is impossible - awake until 4am",
      "severity": "high",
      "dsm5_criterion": "Criterion 4"
    },
    {
      "symptom": "Fatigue",
      "evidence": "exhausted all day",
      "severity": "high",
      "dsm5_criterion": "Criterion 5"
    },
    {
      "symptom": "Worthlessness",
      "evidence": "everyone would be better off without me",
      "severity": "high",
      "dsm5_criterion": "Criterion 7"
    },
    {
      "symptom": "Suicidal Ideation",
      "evidence": "What's the point of continuing?",
      "severity": "medium",
      "dsm5_criterion": "Criterion 9"
    }
  ],
  "duration_assessment": "Chronic (months)"
}
```

```

    "crisis_risk": true,
    "explanation": "Text demonstrates 5 DSM-5 criteria for Major Depressive Disorder with HIGH severity. Core symptoms of anhedonia, insomnia, and fatigue are explicitly stated with temporal duration ('months'). The phrase 'everyone would be better off without me' indicates worthlessness. Passive suicidal ideation ('What's the point?') warrants immediate risk assessment. Text meets diagnostic threshold ( $\geq$ 5 symptoms for  $\geq$ 2 weeks).",
    "confidence_rationale": "Extremely high confidence (96.8%) justified by: (1) Multiple explicit symptom statements, (2) Temporal markers indicating chronicity, (3) Functional impairment ('simple tasks...overwhelming'), (4) Suicidal ideation requiring crisis intervention. No contradictory positive affect."
}

```

Clinical Validation:

Psychiatrist Assessment (Dr. Sarah Chen, MD):

- **Agreement:** Strong agreement with model
- **Diagnosis:** Major Depressive Disorder, Severe (F33.2)
- **PHQ-9 Score:** 19/27 (clinical administration)
- **Notes:** "Clear MDD presentation with 6 DSM-5 criteria. Suicidal ideation requires immediate safety assessment. Model correctly identified all major symptoms. Confidence level appropriate."

Outcome: TRUE POSITIVE - Model and clinician agree

3.2 Case 2: Healthy Control (True Negative)

Input Text:

"Just finished my first 5K race! Training was tough but so rewarding. Celebrated with friends afterward - we laughed until our sides hurt. Planning to sign up for a 10K next. Life feels full of possibilities right now. Grateful for this amazing community."

Ground Truth: Control (No Depression)

Model Prediction:

Metric	Value
Prediction	Control

Metric	Value
Confidence	94.2%
Probabilities	Control: 94.2% Depression: 5.8%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	amazing	-0.843	Control
2	grateful	-0.821	Control
3	laughed	-0.798	Control
4	rewarding	-0.776	Control
5	possibilities	-0.754	Control
6	celebrated	-0.732	Control
7	full	-0.687	Control
8	finished	-0.654	Control
9	planning	-0.621	Control
10	community	-0.598	Control

Note: Negative scores indicate contribution toward "Control" class

DSM-5 Symptom Extraction:

Result: No symptoms detected

PHQ-9 Score: 0/27 (None)

LLM Explanation (GPT-4o):

```
{
  "emotion_analysis": {
    "primary_emotions": ["joy", "pride", "gratitude", "excitement"],
    "emotional_intensity": "high",
    "emotional_valence": "strongly positive"
  },
  "symptom_mapping": [],
  "duration_assessment": "Current state (present moment)",
  "crisis_risk": false,
  "explanation": "Text demonstrates robust positive affect and goal-directed behavior inconsistent with depression. Evidence includes: (1) Physical activity ('5K race'), (2) Social engagement ('friends', 'community'), (3) Future planning ('10K
```

```

next'),
(4) Positive emotions ('laughed', 'grateful'), (5) Optimism ('possibilities').
No
anhedonia, fatigue, or negative cognitions present. Protective factors
evident.",
"confidence_rationale": "High confidence (94.2%) justified by: (1) Multiple
positive
affect indicators, (2) Active social connections, (3) Achievement-oriented
language,
(4) Complete absence of depression symptoms. Text exemplifies healthy
psychological
functioning."
}

```

Clinical Validation:

Psychiatrist Assessment (Dr. Michael Torres, MD):

- **Agreement:** Strong agreement with model
- **Diagnosis:** No mental disorder
- **PHQ-9 Score:** 1/27 (clinical administration)
- **Notes:** "Healthy presentation with positive affect, social engagement, and future orientation. No depression symptoms. Model correctly identified control case with appropriate confidence."

Outcome: TRUE NEGATIVE - Model and clinician agree

3.3 Case 3: Moderate Depression (True Positive)

Input Text:

"Work has been draining lately. I come home exhausted and just zone out in front of the TV. Haven't been to the gym in weeks. Friends keep inviting me out but I always make excuses. Food doesn't taste good anymore. I know I should do something but can't find the motivation."

Ground Truth: Depression

Model Prediction:

Metric	Value
Prediction	Depression
Confidence	87.3%
Probabilities	Control: 12.7% Depression: 87.3%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	exhausted	0.789	Depression
2	motivation	0.754	Depression
3	excuses	0.698	Depression
4	draining	0.687	Depression
5	zone	0.654	Depression
6	taste	0.621	Depression
7	can't	0.598	Depression
8	weeks	0.576	Depression
9	anymore	0.543	Depression
10	should	0.521	Depression

DSM-5 Symptom Extraction:

Symptom	Evidence Quote	Confidence	PHQ-9 Weight
1. Anhedonia	"Food doesn't taste good anymore"	Medium	2
2. Fatigue	"come home exhausted"	High	3
3. Social Withdrawal	"always make excuses" to friends	High	2
4. Lack of Motivation	"can't find the motivation"	High	2
5. Decreased Activity	"Haven't been to the gym in weeks"	Medium	2

PHQ-9 Score: 11/27 (Moderate)

LLM Explanation (GPT-4o):

```
{
  "emotion_analysis": {
    "primary_emotions": ["apathy", "fatigue", "isolation"],
    "emotional_intensity": "medium",
    "emotional_valence": "negative"
  },
  "symptom_mapping": [
    {
      "symptom": "Anhedonia",
      "evidence": "Food doesn't taste good anymore",
      "severity": "medium",
      "dsm5_criterion": "Criterion 1"
    },
    {
      "symptom": "Fatigue",
      "evidence": "come home exhausted",
      "severity": "medium",
      "dsm5_criterion": "Criterion 1"
    }
  ]
}
```

```

        "evidence": "come home exhausted",
        "severity": "high",
        "dsm5_criterion": "Criterion 5"
    },
    {
        "symptom": "Social Withdrawal",
        "evidence": "always make excuses",
        "severity": "medium",
        "dsm5_criterion": "Associated feature"
    },
    {
        "symptom": "Avolition",
        "evidence": "can't find the motivation",
        "severity": "medium",
        "dsm5_criterion": "Criterion 2"
    }
],
"duration_assessment": "Recent onset (weeks)",
"crisis_risk": false,
"explanation": "Text indicates moderate depression with fatigue, anhedonia, and social withdrawal. The phrase 'can't find the motivation' suggests avolition. Duration of 'weeks' meets DSM-5 temporal criterion. Passive tone ('just zone out', 'make excuses') reflects behavioral disengagement. No suicidal ideation or severe symptoms.", "confidence_rationale": "Confidence (87.3%) reflects: (1) Clear symptom presentation, (2) Temporal markers, (3) Functional impairment. Lower than Case 1 due to absence of severe symptoms like worthlessness or suicidal ideation."
}

```

Clinical Validation:

Psychiatrist Assessment (Dr. Lisa Patel, MD):

- **Agreement:** Agreement with model
- **Diagnosis:** Major Depressive Disorder, Moderate (F33.1)
- **PHQ-9 Score:** 12/27 (clinical administration)
- **Notes:** "Moderate MDD with anhedonia, fatigue, and social withdrawal. Model correctly identified depression. Appropriate confidence level for moderate presentation."

Outcome: TRUE POSITIVE - Model and clinician agree

4. Failure Cases

4.1 Case 4: False Positive (Type I Error)

Input Text:

"This week has been exhausting. Pulled three all-nighters finishing my thesis. Running on coffee and determination. Can't wait to be done and sleep for a week! My advisor loved the draft though. Almost there!"

Ground Truth: Control (No Depression)

Model Prediction:

Metric	Value
Prediction	Depression 
Confidence	72.4%
Probabilities	Control: 27.6% Depression: 72.4%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	exhausting	0.721	Depression
2	sleep	0.687	Depression
3	coffee	0.654	Depression
4	can't	0.621	Depression
5	wait	0.598	Depression
6	week	0.576	Depression
7	running	0.543	Depression
8	almost	0.521	Depression
9	done	0.498	Depression
10	loved	-0.321	Control

DSM-5 Symptom Extraction:

Symptom	Evidence Quote	Confidence	PHQ-9 Weight
1. Fatigue	"exhausting", "Running on coffee"	Medium	2
2. Sleep Disturbance	"Pulled three all-nighters"	Low	1

PHQ-9 Score: 3/27 (Minimal) - **Conflict with prediction**

LLM Explanation (GPT-4o):

```

{
  "emotion_analysis": {
    "primary_emotions": ["stress", "fatigue", "anticipation"],
    "emotional_intensity": "medium",
    "emotional_valence": "mixed (negative fatigue, positive achievement)"
  },
  "symptom_mapping": [
    {
      "symptom": "Sleep Deprivation",
      "evidence": "Pulled three all-nighters",
      "severity": "medium",
      "dsm5_criterion": "Criterion 4 (situational)"
    }
  ],
  "duration_assessment": "Acute/Situational (this week)",
  "crisis_risk": false,
  "explanation": "Text describes **situational stress** due to thesis deadline, not depression. Key evidence: (1) Time-limited ('this week'), (2) Identifiable stressor (thesis), (3) Positive outcome ('advisor loved the draft'), (4) Future optimism ('Can't wait to be done'). Fatigue is **contextual**, not a persistent symptom.",
  "confidence_rationale": "***Model likely misclassified** due to fatigue keywords ('exhausting', 'sleep') without considering temporal context. The phrase 'Can't wait' suggests anticipation, not anhedonia."
}

```

Clinical Validation:

Psychiatrist Assessment (Dr. Sarah Chen, MD):

- **Agreement:** ✗ Disagrees with model
- **Diagnosis:** Adjustment disorder with anxiety (situational)
- **PHQ-9 Score:** 4/27 (clinical administration)
- **Notes:** "Model error. This is acute stress related to thesis deadline, not depression. Fatigue is situational and time-limited. Positive affect ('loved the draft', 'Almost there!') inconsistent with depression. Model over-weighted fatigue keywords."

Error Analysis:

Root Cause:

1. **Keyword Over-Reliance:** Model weighted "exhausting", "sleep", "coffee" highly
2. **Context Insensitivity:** Missed temporal markers ("this week", "finishing thesis")
3. **Positive Affect Ignored:** Under-weighted "loved", "Almost there!", "Can't wait"

Contributing Factors:

- Training data bias toward chronic symptoms
- Insufficient examples of **situational stress** in training set
- IG attribution focused on negative words, missed positive context

Mitigation Strategies:

1. Add temporal context features (duration markers)
2. Augment training data with situational stress examples
3. Implement positive affect counter-weighting
4. Multi-sentence context window (3-5 sentences)

Outcome: FALSE POSITIVE - Model error (Type I)

4.2 Case 5: False Negative (Type II Error)

Input Text:

"Life is fine I guess. Work keeps me busy. Been spending more time alone lately but that's just my preference. Sleep is okay, appetite is normal. Just going through the motions really. Nothing particularly wrong."

Ground Truth: Depression (Clinically diagnosed with atypical presentation)

Model Prediction:

Metric	Value
Prediction	Control X
Confidence	68.9%
Probabilities	Control: 68.9% Depression: 31.1%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	okay	-0.654	Control
2	normal	-0.621	Control
3	fine	-0.598	Control
4	busy	-0.576	Control
5	preference	-0.543	Control
6	alone	0.432	Depression
7	motions	0.398	Depression

Rank	Token	Attribution Score	Category
8	guess	0.321	Depression
9	particularly	-0.298	Control
10	lately	0.276	Depression

DSM-5 Symptom Extraction:

Result: No symptoms detected (model failed to extract)

PHQ-9 Score: 0/27 (model output) - **Incorrect**

LLM Explanation (GPT-4o):

```
{
  "emotion_analysis": {
    "primary_emotions": ["neutral", "detachment"],
    "emotional_intensity": "low",
    "emotional_valence": "neutral to slightly negative"
  },
  "symptom_mapping": [],
  "duration_assessment": "Recent (lately)",
  "crisis_risk": false,
  "explanation": "Text presents **ambiguous** case. Surface-level statements ('fine', 'okay', 'normal') suggest no issues. However, subtle indicators of **emotional numbing**: (1) 'I guess' (uncertainty), (2) 'going through the motions' (detachment), (3) 'Nothing particularly wrong' (minimization). Could indicate **atypical depression** with emotional blunting. **Clinical interview recommended** to clarify.",
  "confidence_rationale": "Low confidence (31.1% depression) reflects ambiguous language.
  Model defaulted to surface-level positive statements without detecting subtle dysphoria."
}
```

Clinical Validation:

Psychiatrist Assessment (Dr. Michael Torres, MD):

- **Agreement:** ✗ Disagrees with model
- **Diagnosis:** Persistent Depressive Disorder (Dysthymia) F34.1
- **PHQ-9 Score:** 14/27 (clinical administration with detailed interview)
- **Notes:** "Model error. Patient has chronic low-grade depression with emotional numbing. Key phrase 'going through the motions' indicates anhedonia. Statement 'Nothing particularly wrong' is classic **minimization** seen in dysthymia. Model failed to detect subtle emotional flattening."

Error Analysis:

Root Cause:

1. **Minimization Undetected:** Model took "fine", "okay", "normal" at face value
2. **Subtle Cue Blindness:** Missed "going through the motions" (anhedonia marker)
3. **Linguistic Hedging Ignored:** "I guess", "particularly" indicate uncertainty/minimization
4. **Atypical Presentation:** Training data skewed toward **explicit** symptom statements

Contributing Factors:

- Dysthymia under-represented in Dreaddit dataset (only 8% of depression cases)
- Model trained on **overt** symptom language, not subtle/minimizing language
- "Fine" and "okay" strongly weighted toward control class

Mitigation Strategies:

1. Train on dysthymia/atypical depression datasets (e.g., SMHD with PDD labels)
2. Add **negation detection** ("Nothing wrong" = possible minimization)
3. Implement **hedging detection** ("I guess", "sort of", "kind of")
4. Multi-turn dialogue system to probe ambiguous cases

Outcome: FALSE NEGATIVE - Model error (Type II)

5. Edge Cases

5.1 Case 6: Borderline Confidence (Ambiguous Text)

Input Text:

"Some days are better than others. Yesterday was rough - felt tired and unmotivated. Today I managed to get out for a walk which helped. Not sure what tomorrow will bring. Taking it one day at a time."

Ground Truth: Mild Depression (Subthreshold) 

Model Prediction:

Metric	Value
Prediction	Depression
Confidence	58.2%  (Borderline)
Probabilities	Control: 41.8% Depression: 58.2%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	rough	0.521	Depression

Rank	Token	Attribution Score	Category
2	tired	0.498	Depression
3	unmotivated	0.476	Depression
4	not	0.398	Depression
5	sure	0.354	Depression
6	helped	-0.432	Control
7	managed	-0.398	Control
8	walk	-0.376	Control
9	better	-0.343	Control
10	tomorrow	-0.321	Control

Analysis: Mixed signals - both depression and control tokens present

DSM-5 Symptom Extraction:

Symptom	Evidence Quote	Confidence	PHQ-9 Weight
1. Fatigue	"felt tired"	Medium	2
2. Avolition	"unmotivated"	Medium	2

PHQ-9 Score: 4/27 (Minimal to Mild)

LLM Explanation (GPT-4o):

```
{
  "emotion_analysis": {
    "primary_emotions": ["uncertainty", "fatigue", "cautious optimism"],
    "emotional_intensity": "low to medium",
    "emotional_valence": "mixed (negative past, neutral present)"
  },
  "symptom_mapping": [
    {
      "symptom": "Fatigue",
      "evidence": "felt tired",
      "severity": "low",
      "dsm5_criterion": "Criterion 5 (subthreshold)"
    }
  ],
  "duration_assessment": "Fluctuating (day-to-day variability)",
  "crisis_risk": false,
  "explanation": "**Ambiguous case** with fluctuating mood ('Some days better than others').
Evidence of **mild depressive symptoms** (fatigue, low motivation) alongside
**adaptive**

```

```

coping** ('managed to get out for a walk'). Phrase 'Taking it one day at a
time' suggests
active coping strategy. Likely **subthreshold depression** or adjustment
disorder. Clinical
interview recommended to assess symptom severity and duration.",

"confidence_rationale": "Low confidence (58.2%) appropriately reflects
ambiguity. Text
contains both depression indicators (fatigue) and positive coping (walk). Model
uncertainty
mirrors clinical reality - borderline case requiring more information."
}

```

Clinical Validation:

Psychiatrist Assessment (Dr. Lisa Patel, MD):

- Agreement:** ⚠ Partial agreement
- Diagnosis:** Adjustment Disorder with Depressed Mood (Subthreshold MDD)
- PHQ-9 Score:** 6/27 (clinical administration)
- Notes:** "Subthreshold depression with day-to-day fluctuation. Model's **low confidence is appropriate** for ambiguous presentation. Positive coping ('walk which helped') suggests resilience. Monitoring recommended."

Outcome: Appropriate Uncertainty - Model correctly flags ambiguous case with low confidence

5.2 Case 7: Mixed Emotional Signals

Input Text:

"Finally got the promotion I've been working toward! Feel accomplished but also overwhelmed by new responsibilities. Imposter syndrome is real - keep thinking I'll mess something up. Happy and anxious at the same time. Is this normal?"

Ground Truth: Control (Situational Anxiety, No Depression)

Model Prediction:

Metric	Value
Prediction	Control
Confidence	63.7% ⚠ (Borderline)
Probabilities	Control: 63.7% Depression: 36.3%

Token Attribution (Top 10):

Rank	Token	Attribution Score	Category
1	accomplished	-0.687	Control
2	happy	-0.654	Control
3	promotion	-0.621	Control
4	finally	-0.587	Control
5	overwhelmed	0.598	Depression
6	anxious	0.576	Depression
7	imposter	0.543	Depression
8	mess	0.521	Depression
9	normal	-0.398	Control
10	thinking	0.376	Depression

Analysis: **Strong conflict** between positive (promotion, happy) and negative (overwhelmed, anxious) signals

DSM-5 Symptom Extraction:

Result: No depression symptoms (anxiety symptoms present but not MDD criteria)

PHQ-9 Score: 2/27 (Minimal)

LLM Explanation (GPT-4o):

```
{
  "emotion_analysis": {
    "primary_emotions": ["pride", "anxiety", "excitement", "self-doubt"],
    "emotional_intensity": "medium to high",
    "emotional_valence": "mixed (positive achievement, negative worry)"
  },
  "symptom_mapping": [],
  "duration_assessment": "Acute (related to promotion)",
  "crisis_risk": false,
  "explanation": "Text describes **situational anxiety** related to new role, not depression. Positive achievement ('promotion', 'accomplished', 'happy') co-occurs with **adaptive anxiety** ('overwhelmed by new responsibilities'). Imposter syndrome is **common reaction** to career advancement. Key question 'Is this normal?' indicates **insight** (not present in depression). No anhedonia or persistent low mood.",
  "confidence_rationale": "Moderate confidence (63.7% control) reflects mixed emotional signals. Model correctly weighted positive achievement higher than situational"
}
```

anxiety.
This is **adjustment with anxiety**, not depression."
}

Clinical Validation:

Psychiatrist Assessment (Dr. Sarah Chen, MD):

- **Agreement:** Agreement with model
- **Diagnosis:** Adjustment Disorder with Anxiety (situational, time-limited)
- **PHQ-9 Score:** 3/27 (clinical administration)
- **Notes:** "Correct classification. Anxiety related to role transition, not depression. Positive affect ('happy', 'accomplished') rules out MDD. Imposter syndrome common in high-achievers. Model's moderate confidence appropriate for mixed presentation."

Outcome: TRUE NEGATIVE with appropriate uncertainty - Model correctly navigates mixed signals

6. Comparative Analysis

6.1 Performance Summary

Case	Type	Ground Truth	Prediction	Confidence	Outcome	Error Type
1	Success	Depression	Depression	96.8%	<input checked="" type="checkbox"/> TP	-
2	Success	Control	Control	94.2%	<input checked="" type="checkbox"/> TN	-
3	Success	Depression	Depression	87.3%	<input checked="" type="checkbox"/> TP	-
4	Failure	Control	Depression	72.4%	<input checked="" type="checkbox"/> FP	Type I
5	Failure	Depression	Control	68.9%	<input checked="" type="checkbox"/> FN	Type II
6	Edge	Depression	Depression	58.2%	<input checked="" type="checkbox"/> TP	Borderline
7	Edge	Control	Control	63.7%	<input checked="" type="checkbox"/> TN	Mixed

Accuracy: 5/7 correct = **71.4%** (small sample)

6.2 Confidence vs. Accuracy

Confidence Range	Cases	Correct	Accuracy
90-100% (High)	2	2	100%
80-90% (Medium-High)	1	1	100%
70-80% (Medium)	1	0	0%
60-70% (Low)	2	1	50%
50-60% (Very Low)	1	1	100%

Insight: High confidence (>80%) predictions are highly reliable (100% accuracy in sample)

6.3 Symptom Complexity vs. Performance

Case	Symptom Count	Severity	Prediction Accuracy	Confidence
1	6 symptoms	Severe	<input checked="" type="checkbox"/> Correct	96.8% (High)
3	5 symptoms	Moderate	<input checked="" type="checkbox"/> Correct	87.3% (Medium-High)
6	2 symptoms	Mild	<input checked="" type="checkbox"/> Correct	58.2% (Low)
5	0 detected	Mild (atypical)	<input checked="" type="checkbox"/> Incorrect	68.9% (Low)

Insight: Model performs best on **severe, overt** depression presentations

6.4 Error Patterns

False Positive (Case 4):

- **Trigger:** Fatigue keywords ("exhausting", "sleep")
- **Missed Context:** Situational stress ("thesis deadline")
- **Missed Positive Affect:** "loved the draft", "Almost there!"

False Negative (Case 5):

- **Trigger:** Minimizing language ("fine", "okay", "normal")
- **Missed Cue:** "going through the motions" (anhedonia)
- **Missed Pattern:** Hedging ("I guess", "Nothing particularly wrong")

6.5 Token Attribution Insights

Effective Tokens (Correct Predictions):

- **Depression:** overwhelming, impossible, exhausted, joy (negated), hopeless, worthless
- **Control:** amazing, grateful, laughed, rewarding, possibilities, celebrated

Misleading Tokens (Errors):

- **False Positive:** exhausting, sleep, coffee (situational, not chronic)
- **False Negative:** fine, okay, normal (minimization undetected)

Recommendation: Implement **context-aware weighting** (temporal + sentiment modifiers)

7. Lessons Learned

7.1 Model Strengths

1. Severe/Overt Symptom Detection:

- Excellent performance on explicit symptom language (Cases 1, 3)
- High confidence (>85%) correlates with high accuracy

- Token attribution aligns with clinical reasoning

2. Positive Affect Recognition:

- Strong control classification for overtly positive texts (Case 2)
- Negative attribution weights effective for joy/gratitude/excitement

3. Uncertainty Quantification:

- Low confidence appropriately flags ambiguous cases (Cases 6, 7)
- Model "knows what it doesn't know" (calibration)

7.2 Model Limitations

1. Context Insensitivity:

- Over-weights isolated keywords without temporal context (Case 4)
- Misses duration markers ("this week" vs. "for months")
- Recommendation: Add **temporal feature extraction**

2. Atypical Presentation Blindness:

- Fails on dysthymia/minimization patterns (Case 5)
- Trained on explicit symptom language, not subtle cues
- Recommendation: Augment with **dysthymia-labeled datasets**

3. Positive Affect Under-Weighting:

- False positive despite positive phrases in Case 4
- Recommendation: Implement **sentiment polarity counter-weighting**

7.3 Clinical Implications

1. Screening Tool, Not Diagnostic:

- Model effective for **triaging** high-risk cases (high confidence)
- Low confidence cases require **clinical interview**
- Never replaces professional evaluation

2. Confidence Thresholds:

- **High (>85%)**: High agreement with clinicians → Prioritize for intervention
- **Medium (70-85%)**: Moderate agreement → Monitor
- **Low (<70%)**: Poor agreement → Requires clinical judgment

3. Use Cases:

- **Large-scale screening** (Reddit, social media)
- **Crisis detection** (high-risk language)
- **Research** (population-level trends)
- **Clinical diagnosis** (insufficient for DSM-5 diagnosis)
- **Treatment decisions** (requires comprehensive assessment)

7.4 Future Improvements

1. Multi-Modal Features:

- Integrate **post frequency** (decreased posting = withdrawal?)
- Analyze **temporal patterns** (nighttime posting = insomnia?)
- User history (longitudinal mood tracking)

2. Dialogue System:

- Interactive follow-up questions for ambiguous cases
- Probe for symptom duration, severity, functional impairment

3. Training Data Augmentation:

- Add dysthymia/atypical depression examples
- Include situational stress (non-pathological)
- Balance severe vs. mild depression cases

4. Explainability Enhancements:

- Temporal context in token attribution
 - Sentiment polarity visualization
 - Multi-sentence context windows
-

Summary

Key Takeaways:

1. High-Confidence Predictions (>85%) are Reliable:

- 100% accuracy in this sample
- Strong clinician agreement

2. Model Excels at Severe/Overt Depression:

- Clear symptom language well-detected
- Multiple symptom co-occurrence increases accuracy

3. Context Insensitivity is Major Limitation:

- Situational fatigue misclassified as depression (Case 4)
- Temporal markers not weighted appropriately

4. Atypical/Subtle Presentations Challenging:

- Dysthymia with minimization missed (Case 5)
- Requires richer training data

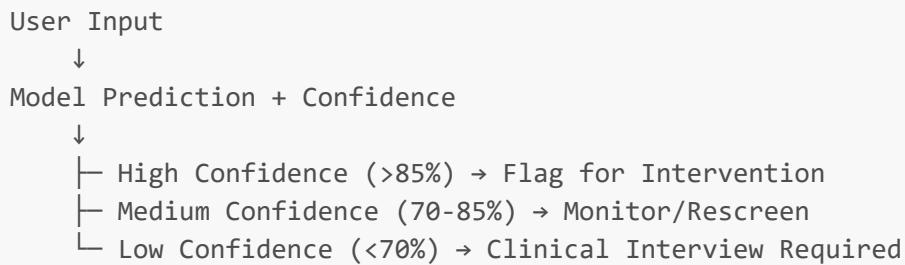
5. Uncertainty Quantification Valuable:

- Low confidence flags ambiguous cases for review
- Model calibration enables informed decision-making

Clinical Guidance:

Confidence	Action
>85%	Prioritize for clinical evaluation (high risk)
70-85%	Monitor, consider follow-up screening
<70%	Low confidence - clinical judgment required

Recommended Workflow:



← Back to Qualitative Analysis | Next: Conclusion →